

quick-commerce-rag

LangChain Pipeline

A retrieval-augmented chatbot for packaged-food ingredient, nutrition, allergen, and FSSAI regulatory questions — hybrid retrieval, LLM tool-calling routing, typed-claim generation, and post-generation groundedness verification, built end to end on LangChain.

Hybrid Retrieval (BM25 + Dense)

Cross-Encoder Reranking

LLM Tool-Calling Router

Groundedness Verification

Conversation Memory

RAGAS-style Evaluation

CONTENTS

1. What this is

2. Data

3. Request flow

4. Hybrid retrieval

5. Generation and groundedness

6. Conversation memory

7. LangChain integration

8. Evaluation

9. Application-layer guardrails

10. Serving layer

11. Deployment

12. Design principles this project actually follows

A deep dive into how this project works end to end: the problem it solves, the data it's built on, how a question travels through the system, and the design decisions behind each piece — including the ones that were tried and reverted.

This document describes the LangChain-native pipeline only. Everything here is implemented with LangChain primitives (`ChatGroq` , `ChatHuggingFace` , LCEL chains, a custom `BaseRetriever`) rather than hand-rolled HTTP calls.

1. What this is

A RAG chatbot embedded in a mock quick-commerce app UI. It answers questions about 23 real packaged food products — ingredients, nutrition, allergens, and FSSAI (India's food regulator) compliance — by combining:

- **Deterministic structured data** (SQLite) for facts that have one correct answer: calories, ingredient lists, allergen declarations, pack size.
- **Retrieval-augmented generation** over a curated knowledge base (FSSAI additive/labelling rules, WHO nutrition guidance, an ingredient encyclopedia) for questions that need interpretation or regulatory context: "is this preservative safe?", "can this product claim to be a good source of protein?"

The goal wasn't to prove an LLM can answer these questions — a raw LLM call already can, badly, by hallucinating regulatory limits it was never given. The goal was building the retrieval, grounding, and verification machinery that makes those answers *trustworthy*: every claim is typed and cited, and a post-generation check verifies a cited number actually appears in the source it claims to come from.

2. Data

2.1 Structured store — SQLite (`db/products.sqlite`)

One row per product, loaded from `data/raw/products_compiled.json` by `ingestion/load_products.py`. Nested fields (nutrition, allergens, pack size) are stored as JSON text columns and queried with `json_extract()` — deliberately not normalized into separate tables, since the dataset is only 23 products and normalization would add joins without adding value at this scale.

Fact-lookup questions ("how many calories does this have?") bypass retrieval and generation entirely — they're answered by a direct SQL query, with zero LLM calls. This is a real cost optimization: a large fraction of real questions about a product are simple fact lookups, and routing those around the LLM entirely means faster answers with no hallucination surface at all for that class of question.

2.2 Vector store — Qdrant

One flat collection (`kb_baseline`), holding markdown chunks from four knowledge base files:

File	Content	Chunks
<code>fssai_knowledge_base.md</code>	Additive limits, labelling rules, marketing-claim eligibility	52
<code>nutrition_knowledge_base.md</code>	WHO Sugars 2015 + Healthy Diet guidance	15
<code>ingredient_knowledge_base.md</code>	Tier 1: INS-coded additives (aspartame, preservatives, colours...)	41
<code>ingredient_kb_tier2.md</code>	Tier 2: base/whole ingredients (wheat, milk, sugar...)	12


Chunked by `ingestion/parse_kb.py` on markdown `##` / `###` headers — no smart table handling, no overlap, no chunk-size tuning. This is intentional: the chunking strategy isn't the thing this project set out to demonstrate, so it stays simple rather than over-engineered.

Each chunk carries payload metadata when the source markdown declares it: `doc_type`, `entity` (which ingredient/additive it's about), `ins_no` (the regulatory code, which can be a **compound** value like `"627/631/635"` for a multi-code entry — every place this gets compared against a single code splits it first, since a naive string-equality check silently excludes every multi-code entry from retrieval otherwise), and `comparison_group` (an author-assigned tag linking chunks across files that answer a comparison question together, e.g. `sugar_vs_sweetener`).

Embedded with `BAAI/bge-small-en-v1.5`. `bge` models require an instruction prefix on the **query** side only, never on documents — a documented model requirement (`BGE_QUERY_PREFIX` in `config.py`), not a retrieval-strategy choice.

Metadata coverage is genuinely thin and uneven across the KB — not every chunk has a real `doc_type / entity / ins_no`. Every filtering/scoring mechanism downstream is deliberately **inclusive by default**: it only excludes a candidate when it has specific, positive evidence against it (e.g. "this chunk's `ins_no` doesn't match any code this product actually declares"), never on the absence of metadata.

3. Request flow



3.1 The deterministic fast path

`classify_query()` resolves which of the 23 catalog products (if any) a query is about by scoring name/brand token overlap — ties or a zero score return "no match" rather than guessing. If a product resolves **and** the query matches a known fact-field pattern **and** no regulatory-override term ("safe", "legal", "compare", "too much", etc.) is present, it routes straight to SQL with zero LLM involvement.

This fast path is deliberately narrow — a real cost optimization for confident, unambiguous cases only. Everything it doesn't confidently resolve falls through to the LLM tool-calling router below, not to a second keyword table. An earlier version of this project used a second layer of hand-written intent-classification keyword lists for the retrieval side; it was deleted outright after repeatedly hitting the same failure mode — a phrase list that didn't anticipate one new wording — which no amount of list-patching can close in general. LLM tool-calling replaced it because understanding "does this question need the knowledge base" is exactly the kind of open-ended language task an LLM is suited for and a keyword table isn't.

3.2 LLM tool-calling router

A single-round decision (not an open-ended agentic loop — every real routing failure found in development was fixable by one correct tool choice, none needed sequential tool chaining) with `tool_choice="required"` : the model must commit to at least one tool rather than declining. The system prompt instructs it to call a structured tool **and** `search_knowledge_base` together whenever a question needs both the product's own data and general regulatory/nutrition context (e.g. "is this healthy?", "why does this need a preservative?").

If only structured tool(s) fire, their already-cited output returns directly — no second LLM call. This generalizes the fact-lookup route's "instant, zero-LLM-cost" property to the full structured-data surface (comparisons, ingredient/allergen checks) instead of needing its own keyword pattern for each case.

LLM tool-selection has an irreducible error rate — this is treated as a property of the architecture, not a bug to chase to zero with more prompt engineering. Where a systematic gap was found (the model reliably under-calling `search_knowledge_base` for a specific, recurring question shape — evaluative/health-judgment questions, regulatory-limit questions, claim-eligibility questions, dietary/nutritional-verdict questions), the fix was a **structural pattern detector**: a regex that generalizes across phrasing, purely additive (it only ever adds more grounding to an answer, never blocks or overrides the model's own tool choices), never a hand-typed list of exact phrases. Four such detectors exist, covering health judgments, regulatory limits, claim eligibility, and dietary/nutritional verdicts — each was found via either a live user report or a proactive synthetic stress-test batch, and each is verified with regression tests.

4. Hybrid retrieval

When `search_knowledge_base` fires, the query goes through:

1. **BM25 + dense fusion** — BM25 (keyword/lexical) top-20 and dense (semantic/embedding) top-20 candidates, combined by Reciprocal Rank Fusion. Keyword and semantic search catch different failure modes: BM25 finds exact regulatory terms and product names dense search can miss; dense search finds paraphrases and conceptual matches BM25 can't.
2. **Narrow, evidence-based exclusion** — a candidate is dropped only when the query resolved to a specific product **and** the candidate is an ingredient chunk with a specific INS code that isn't among that product's actual declared codes. A functional-class boost (sourced from the knowledge base's own section headings, not a hand-typed list) also *excludes* a genuinely-declared ingredient chunk belonging to a different functional class than the one the query asked about (e.g. a query about "the colour additive" shouldn't be answered by a preservative chunk just because both are declared ingredients).
3. **Cross-encoder reranking** (`cross-encoder/ms-marco-MiniLM-L-6-v2`) of the top ~15 fused candidates, sigmoid-normalized.
4. **Corroborated-chunk trim** — a reranked chunk is kept only if it's independently confident on the cross-encoder alone, OR it's a top-2 pick by *either* the cross-encoder ranking or the raw BM25 ranking. This exists because the cross-encoder is real but demonstrably miscalibrated on this KB's long, multi-topic technical chunks — a genuinely-needed chunk was measured scoring as low as 0.055 while a topically-adjacent-but-wrong chunk scored 0.631. A single fixed score-floor trim was tried first and reverted after it dropped a chunk the model then hallucinated a citation for; the two-signal corroboration check is what actually shipped.
5. **Corrective retry** — if the top reranked score is below a threshold (0.3, validated against real score distributions, not guessed), the query is rewritten and steps 1–4 run once more.
6. **BM25-consensus skip** — the retry is skipped, and the original top-1 result trusted, when that result is *also* a landslide BM25-score winner ($\geq 1.4\times$ the runner-up, both scores real and substantial). This exists because the cross-encoder's absolute score genuinely can't always be trusted (see step 4) — for one specific real question, the correct chunk was the cross-encoder's own #1 pick and a landslide BM25 winner, yet still scored below the retry threshold; the "corrective" retry was actually making a correct answer worse by rewriting away from it.
7. **comparison_group override** — if both the original and retried query still fail to clear the threshold, but the candidate pool contains two or more chunks sharing a real, pre-authored `comparison_group` tag spanning different source files, that pairing is returned instead of an insufficient-evidence result. Deliberately narrow: it only fires on an explicit tag match, not general multi-hop synthesis. A genuinely untagged multi-chunk-synthesis question (no single chunk has the full answer, no tag exists) correctly falls through to insufficient-evidence — general sub-query decomposition for that class of question is a known, explicitly out-of-scope gap.

If nothing clears the bar even after all of the above, the pipeline returns an honest [UNCERTAIN] Insufficient evidence ... message rather than guessing.

5. Generation and groundedness

5.1 Typed-claim generation

The system prompt (`generation/llm.py`) requires every non-trivial claim to be tagged `[FACT]` / `[REGULATORY]` / `[INTERPRETATION]` / `[DERIVED CALCULATION]` / `[UNCERTAIN]` , each with a source citation, and explicitly forbids filling retrieval gaps from the model's own outside knowledge. This matters most exactly where it's hardest to enforce: an additive whose retrieved context is a genuine, confirmed regulatory gap (the KB simply doesn't cover it) needs the model to say so, not to confidently invent a plausible-sounding limit.

5.2 Groundedness check

A post-generation, numeric-only verification: each `[FACT]` / `[REGULATORY]` / `[INTERPRETATION]` claim's cited number (if it has one) is checked against the actual text of the chunk it cites, and annotated with an inline  `[UNVERIFIED - ... ]` marker on mismatch — never silently regenerated or edited.

Deliberately numeric-only, not general NLI. Three general entailment/contradiction-checking designs (whole-chunk entailment, per-sentence contradiction, whole-chunk contradiction) were built and tested against real chunks; all three produced real false positives on this KB's long, multi-topic technical chunks. A claim with no extractable number passes through unchecked — a known, stated limitation, not an oversight. This is a concrete instance of a broader theme in this project: several "obviously more thorough" designs were tried, measured against real data, and reverted in favor of a narrower mechanism that actually works reliably — see also the corrective-retry threshold (§4) and the groundedness scope itself.

6. Conversation memory

Follow-up questions ("is that too much?", "what about the diet version?") are expanded into self-contained queries *before* they reach routing — `conversation/resolve.py`, template-based, not an LLM call. This is a deliberate choice, for two reasons: it needs to keep working even when every configured LLM provider is exhausted (a real, recurring constraint during development — see §8), and template-based resolution is easier to verify and won't casually reintroduce the same product name into an already-scoped query in a way that skews retrieval toward product/ingredient chunks over the generic guidance chunks a question actually needs.

`conversation/state.py` tracks `known_facts` (attribute → value/unit/source, established across turns), the active product, and the active topic. Two checks build on this:

- **Consistency check** (`generation/consistency.py`) — catches a freshly generated answer contradicting an already-established fact (a genuine numeric contradiction) or claiming a known attribute is "unavailable" when it was already given earlier in the conversation. A guideline number appearing near the same unit as the product's own known value (e.g. "≤50g added sugars" next to a product's real "47.4g") is deliberately *not* flagged — several real false-positive shapes of this exact kind were found and fixed via live testing, each locked in with a regression test.
- **Groundedness check** (§5.2) — checks a claim against its *cited chunk*.

These check two different things and are both run: a claim can be perfectly grounded in its cited source while still contradicting something the user was already told two turns ago.

7. LangChain integration

Everything an LLM makes a real decision in — tool-routing and final-answer generation — goes through LangChain-native call paths (`ChatGroq` , `ChatHuggingFace` , an LCEL prompt chain). Everything that *isn't* an LLM call — BM25 fusion, reciprocal rank fusion, cross-encoder reranking, the comparison-group tag match, structured SQL dispatch, conversation-state bookkeeping, the regex safety-nets from §3.2 — is plain Python, reused directly rather than reimplemented "in LangChain style." LangChain has no abstraction for "cross-encoder-rerank a fused RRF pool" or "does this query mention a health condition" — these are just Python, and reimplementing already-tested, hard-won logic a second time would only add a second, divergent copy to keep in sync.

A custom retriever/embeddings wrapper is used instead of `langchain_qdrant` 's default integration where the project's Qdrant payload shape (flat, metadata-scoped fields) doesn't match what that integration expects out of the box.

7.1 Gateway: multi-key rotation, proactive budgeting, fallback

`groq_gateway_invoke()` reimplements, LangChain-natively, a three-layer call strategy:

1. **Proactive budget check** — before attempting a Groq key, a per-key daily token-budget ledger (persisted to disk, survives process restarts) is checked against 95% of Groq's real 200,000 token/day per-key limit, using a request-size estimate. A key predicted to be too close to exhaustion is skipped *before* spending a real network round-trip finding out.
2. **Reactive rate-limit catch** — if a key is attempted and actually rate-limited anyway (quota spent by something outside this ledger's visibility), the next key is tried.
3. **Hugging Face fallback** — once every configured key is proactively-or-actually exhausted, falls back to a `ChatHuggingFace` model. Multiple Groq keys configured means real headroom before this fallback is ever needed in practice.

A real compatibility bug was found and fixed in this fallback path: the installed `langchain_huggingface` version rejects any non-empty `tool_choice` — including "auto" — unless exactly one tool is bound, and this pipeline's routing call always binds four. The fix passes `None` whenever more than one tool is given, letting the fallback model see every tool and decide freely rather than crashing.

A global, disk-persisted **exact-match cache** (LangChain's `SQLiteCache`) sits in front of every model call — an identical (prompt, model, params) tuple returns instantly with zero token spend and zero impact on the budget ledger, since a cache hit never actually reaches the provider.

8. Evaluation

A hand-rolled RAGAS-equivalent harness (`eval/ragas_metrics.py` , `eval/run_ragas_eval.py`) — not the `ragas` package itself, since that depends on LangChain's own LLM-wrapper abstractions for its judge calls in a way that would bypass this project's own quota management. Computes four metrics via a decompose-then-verify LLM-as-judge technique:

- **Faithfulness** — does the generated answer's content hold up against the retrieved context (KB chunks *and* any legitimately-cited structured facts)?
- **Answer relevancy** — does the answer actually address the question asked?
- **Context precision** — how much of what was retrieved was actually relevant?
- **Context recall** — did retrieval surface what a good reference answer would need?

Running this harness against real questions found and drove fixes for several real bugs during development — a regression in a query-rewriting tool schema, an ingredient-matching gap, three routing/tool-selection gaps, and a retrieval-filtering bug (the compound- `ins_no` issue mentioned in §2.2) that had been silently excluding every multi-code KB entry from retrieval. This is the pattern this project follows generally: eval isn't a report generated once at the end, it's a tool that actively surfaces where the system is actually wrong.

Offline regression tests (no LLM/network cost — pure-function checks with fixture data, or LLM calls monkeypatched to cost nothing) exist per module and are run any time that module changes: retrieval fusion/exclusion logic, the consistency false-positive fixes, the token-budget ledger, the security guardrails below, and the RAGAS metrics' own scoring functions.

9. Application-layer guardrails

Scoped to `/api/chat` specifically, since that's the endpoint that spends real provider quota:

- **Rate limiting** — a sliding window, per-client-IP.
- **Query length bound** — enforced by the request schema.
- **Prompt-injection heuristic** — a narrow regex matching only near-unambiguous instruction-override attempts ("ignore all previous instructions", "reveal your system prompt"). Deliberately not a blunt keyword blocklist: this KB's own domain language legitimately contains phrases like "ignore the label" or "disregard the packaging claim" that a blocklist would false-positive on. Flags the response but never hard -blocks — the real structural defense against off-topic or fabricated answers is the typed-claim system prompt (§5.1) and the groundedness check (§5.2), neither of which this heuristic changes.

10. Serving layer

Two FastAPI apps front the same pipeline:

- **The main app** (port 8000) — serves the built React frontend and `/api/chat / /api/products` . Every chat request goes through the full routing → retrieval/tools → generation → groundedness → consistency flow described above. This is the one meant to be deployed publicly.
- **A secondary, API-only app** (port 8001) — no frontend, adds an SSE streaming endpoint and extra observability fields (which tools fired this turn, per-stage latency, real token usage) that the main app doesn't expose. Useful for inspecting what the pipeline actually did on a given question, or for a client that wants to render tokens as they arrive rather than waiting for the full answer.

Both share one cached bundle of loaded models (embedding model, cross -encoder, BM25 index, Qdrant client) — built once at process startup, not reloaded per request.

11. Deployment

Containerized as a multi-stage Docker build: the React frontend is built in a Node stage, then copied into a Python runtime stage alongside the source. `products.sqlite` is built at image-build time (deterministic, offline -derivable data — no network call needed during the build).

Qdrant is **not** embedded in the image — it's a separate, stateful service, pointed at via `QDRANT_URL` / `QDRANT_API_KEY` (a managed instance, e.g. Qdrant Cloud) rather than the bare `localhost:6333` local development uses. The one manual step a deploy needs that the container can't do for itself: running the ingestion script once against that managed instance before the app can answer anything, since embedding/upserting needs to happen exactly once per knowledge-base version, not on every container boot.

12. Design principles this project actually follows

A few threads run through the decisions above, worth naming directly:

- **Prefer the narrow, evidence-based mechanism over the broad, clever one.** The exclusion/boost logic in retrieval, the groundedness check's numeric-only scope, the comparison-group override's exact-tag requirement — each of these is a case where a more general version was either tried and reverted (produced false positives on real data) or deliberately not built (the data didn't support it working reliably).
- **Cost-aware by default, not as an afterthought.** The deterministic fact-lookup fast path, the "structured-tools-only" instant path, the proactive token-budget ledger, and the exact-match LLM cache all exist because every unnecessary LLM call is real latency and real cost at scale — this shows up as an architectural pattern (bypass generation entirely when possible), not just an infrastructure detail.
- **An eval harness that's actually used, not just built.** Every metric in §8 has a real story of a bug it found and a fix it drove — the harness is treated as an active diagnostic tool run against real questions, not a one-time report.
- **Honest failure over confident guessing.** Insufficient retrieved evidence returns `[UNCERTAIN]`, not a plausible-sounding fabrication. A genuine regulatory gap in the source data is reported as a gap, not papered over. This is the single property the typed-claim system prompt and the groundedness check both exist to protect.